

Credit Card Fraud Detection



Nirbhaya Shaji, Lúcia Moreira

Introduction to Data Science Project, FCUP, University of Porto

December 17, 2019

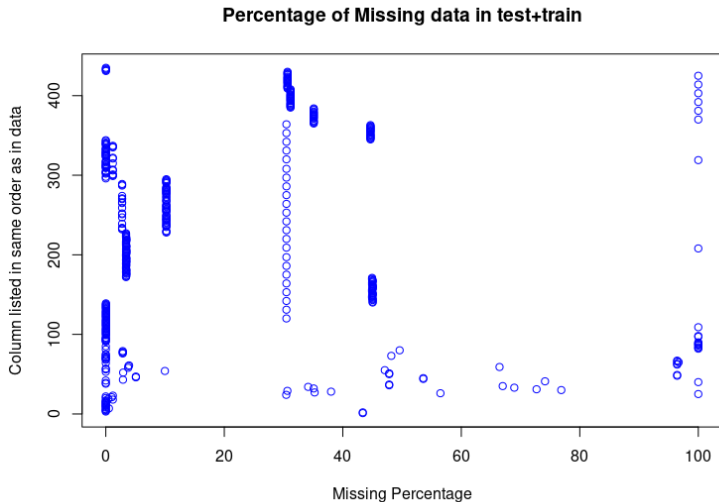
Kaggle: IEEE-CIS Fraud Detection

- Predicting the probability that an online transaction is fraudulent
- Vesta Corporation - e-commerce payment solutions
- Data is broken into two files: identity and transaction.
- Categorical and Numerical features
- train_transaction, identity.csv - the training set
- test_transaction, identity.csv - the test set



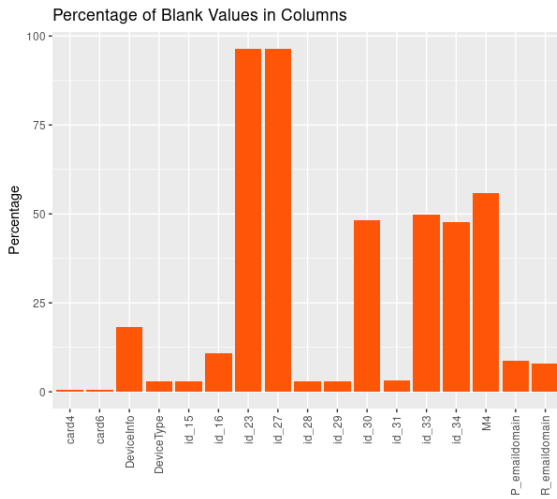
Missing Data

- Transaction and Identity merged over TransactionID
- Test and Train Data Combined

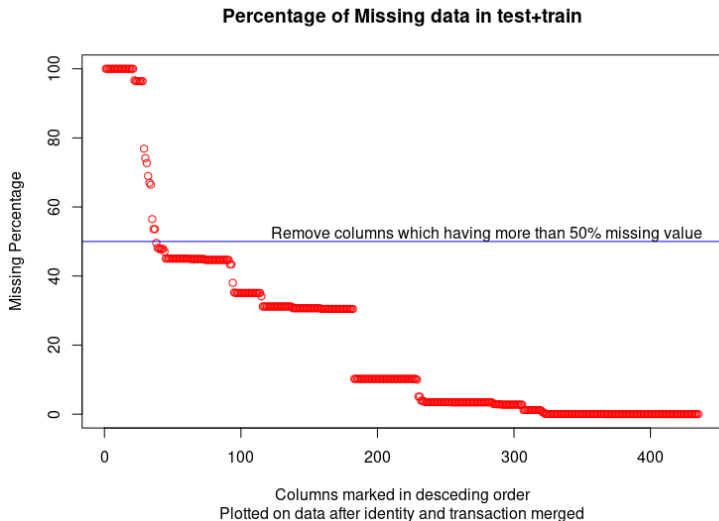


Sparsity

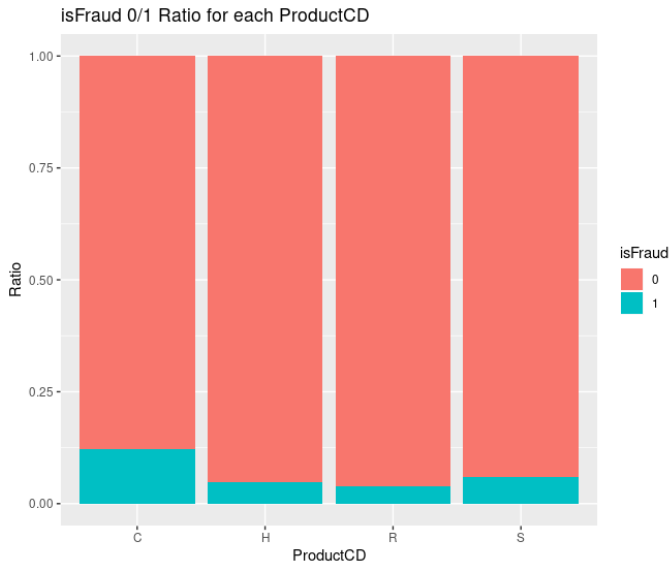
- Random Sparsity
- 15 percent of the data has NA values (17514759/114169860)
- Columns with blank values present



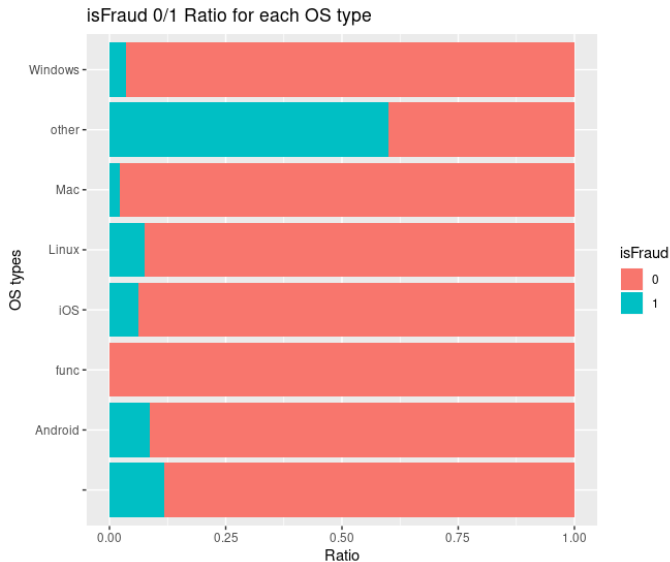
- Columns with more than 50 percentage NA removed for further EDA



ProductCD

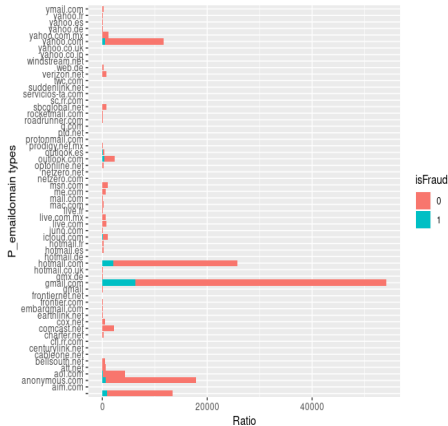


OS type




Email Domain

isFraud 0/1 Ratio for each P_emaildomain




isFraud 0/1 Ratio for each R_emaildomain



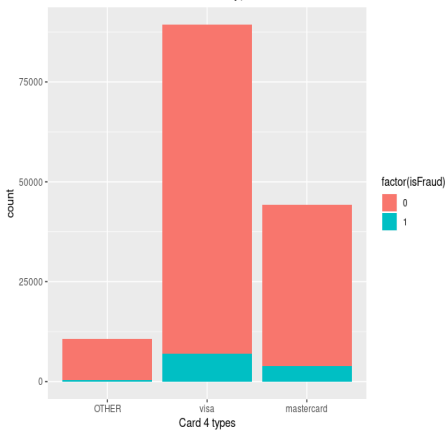
- Pretty much same as P-emaildomain with an exception of scranton.edu

Transaction Amount Box-plot

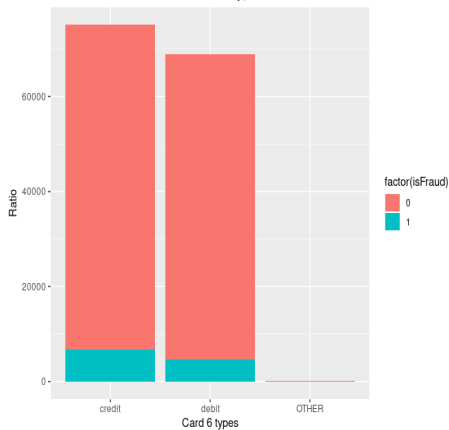


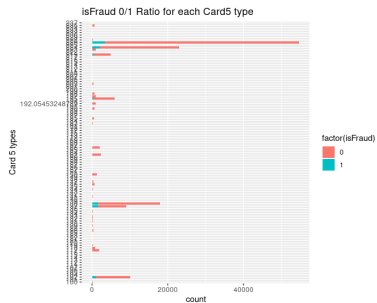
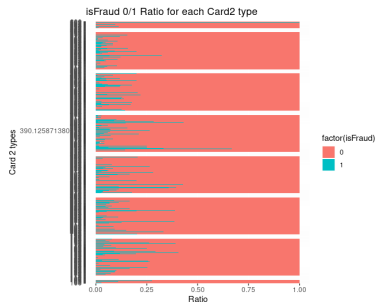
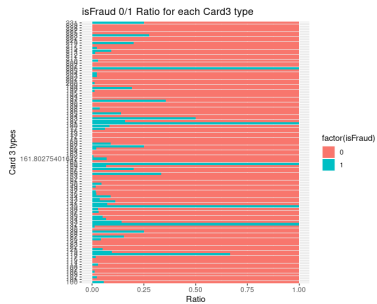
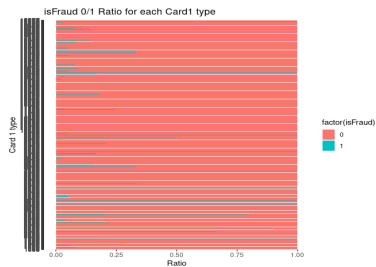
CardType 4: Visa/Master/Other and CardType 6: Credit/Debit/Other

isFraud 0/1 Ratio for each Card4 type



isFraud 0/1 Ratio for each Card6 type



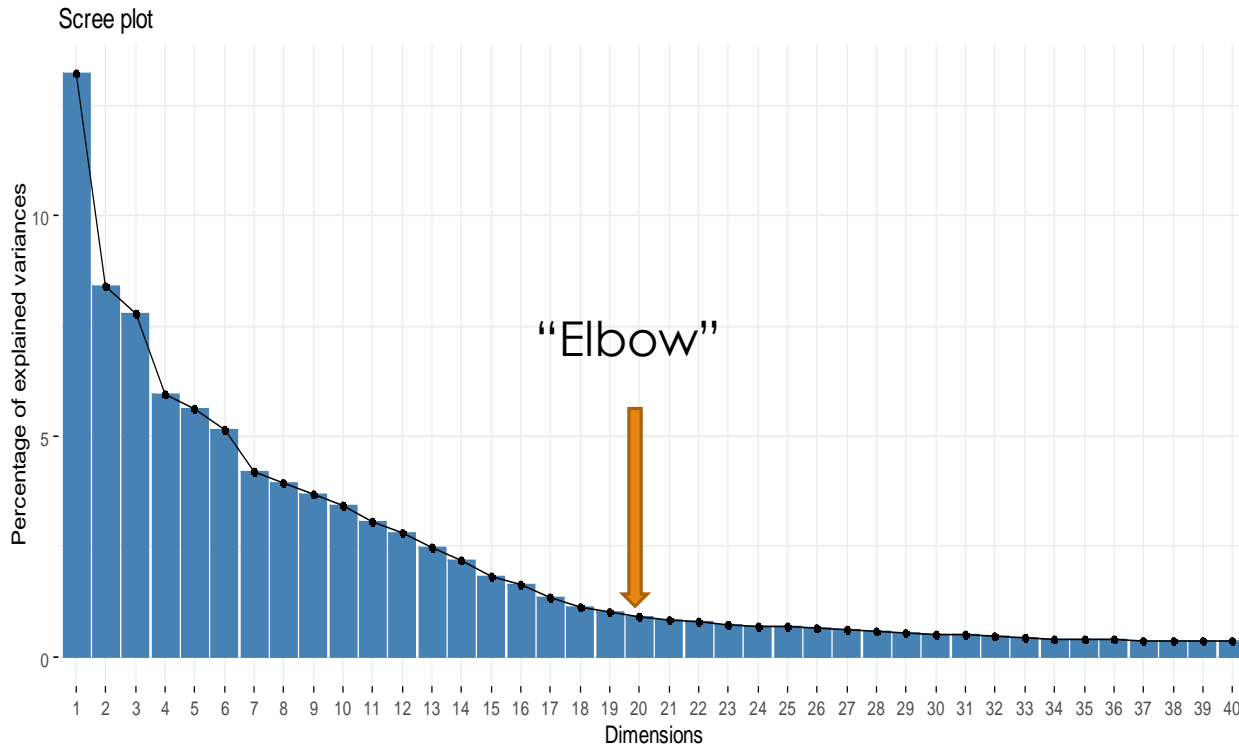


Strategy I – decrease dataset size

- ▶ Very high sparse initial dataset files:
 - ▶ Transactions (590540, 394): **41 %** (670 MB)
 - ▶ Identity (144233, 41): **36 %** (26 MB)
- ▶ Merge by identification ID (144233, 434): **27 %**
 - ▶ Combine information relevant in both files
 - ▶ Also allows decreasing dataset size
- ▶ Drop columns with more than 50% NaNs (144233, 303): **11%**
 - ▶ Drop 131 columns – allowed further decrease in dataset size (180 MB yet)
- ▶ Unbalanced feature was kept:
 - ▶ 8 % fraud
 - ▶ 92% not fraud

Strategy II – Decrease number of numerical variables

- We still have **270 numerical variables** in the new dataset
- **PCA analysis:**
 - Mean was input in the remaining NaNs



Comply with the 3 common criteria



Component	eigenvalue	% of variance	cumulative % of variance
1	34,9	13,2	13,2
2	22,2	8,4	21,6
3	20,5	7,8	29,4
4	15,7	6,0	35,4
5	14,9	5,6	41,0
6	13,6	5,1	46,1
7	11,1	4,2	50,3
8	10,4	4,0	54,3
9	9,7	3,7	58,0
10	9,0	3,4	61,4
11	8,1	3,1	64,4
12	7,4	2,8	67,2
13	6,6	2,5	69,7
14	5,8	2,2	71,9
15	4,8	1,8	73,7
16	4,3	1,6	75,4
17	3,6	1,4	76,7
18	2,9	1,1	77,8
19	2,7	1,0	78,9
20	2,4	0,9	79,8
21	2,2	0,8	80,6



Strategy III – Identify relevance of the categorical variables

- ▶ Data set has 33 categorical variables
- ▶ Random Forest – handles NaNs
- ▶ Ensemble with 500 trees
- ▶ CPU Memory intensive
- ▶ Categorical variables with more than 53 factors: removed!
 - ▶ 15 variables removed

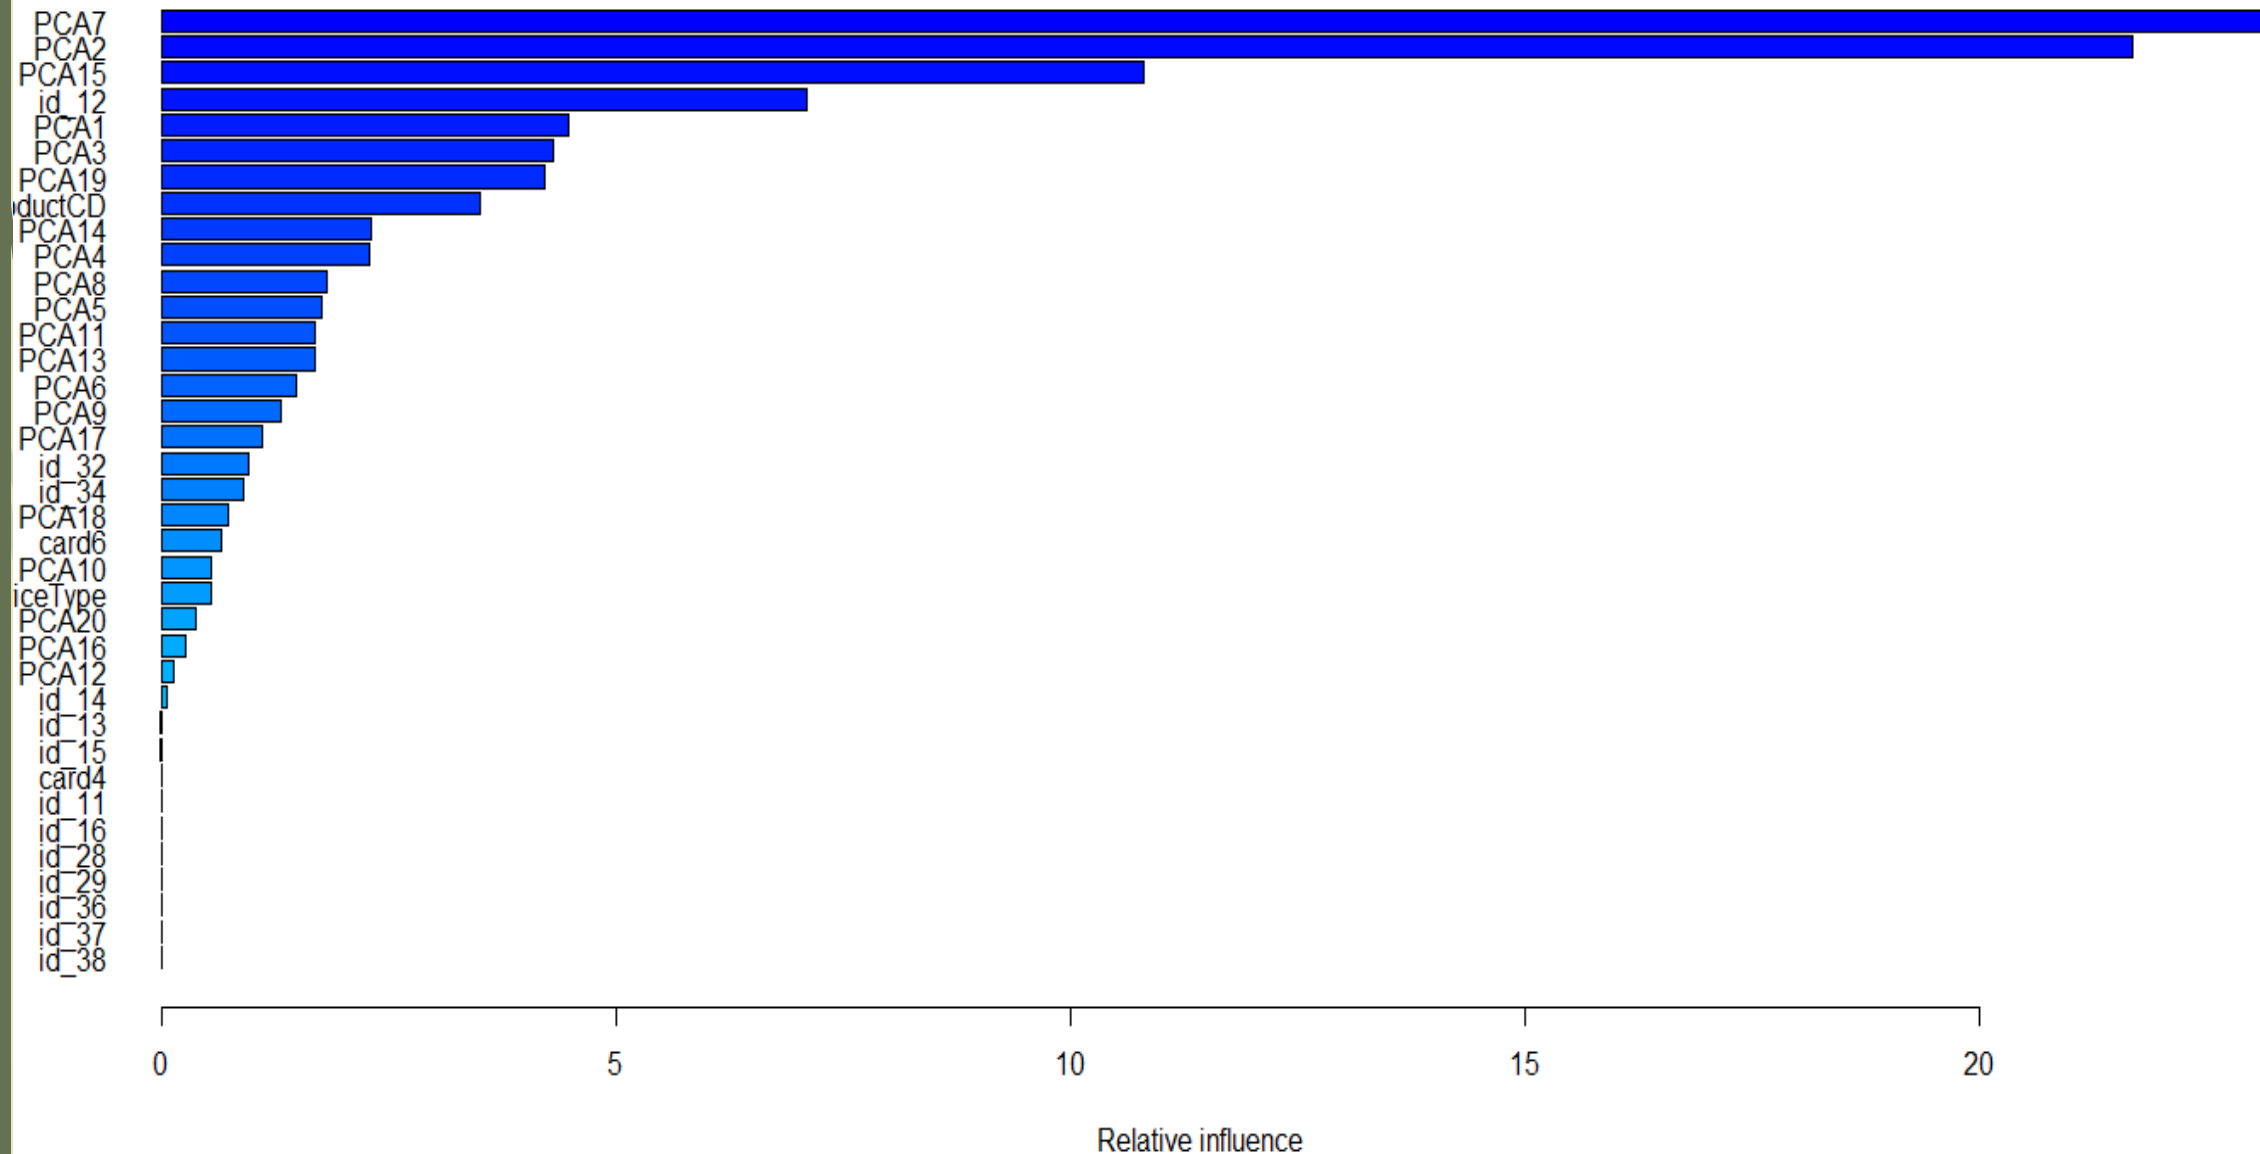
Strategy III – Random Forest



Model – Gradient Boosting Machine

- **Relevant algorithm for unbalanced data sets**
 - **20 numerical dimensions from PCA**
 - **18 categorical significant variables with up to 53 factors**
-
- ▶ Time domain consideration:
 - ▶ 50% past for training
 - ▶ 50% future for testing
 - ▶ Hyperparameters:
 - ▶ Holdout method with 30% future for validation based on accuracy:
 - ▶ Shrinkage (regularization),
 - ▶ nr. trees
 - ▶ depth
-
- ▶ Random splitting
 - ▶ 65 % for training
 - ▶ 35 % for testing
 - ▶ Hyperparameters:
 - ▶ 10-fold CV based on accuracy:
 - ▶ shrinkage (regularization)
 - ▶ nr. trees
 - ▶ depth

Model – Gradient Boosting Machine



Model – Gradient Boosting Machine

- Time domain consideration
 - Logistic Regression (Bernoulli distribution):
 - Accuracy train: **96.2 %**

	0	1
prediction		
0	68588	2526
1	253	1133

- Accuracy test: **89.8 %**

	0	1
prediction		
0	63696	6945
1	377	714

- AdaBoost (exponential loss for 0-1 outcomes):
 - Accuracy train: **96.0 %**

	0	1
prediction		
0	68595	2626
1	246	1033

- Accuracy test: **89.7 %**

	0	1
prediction		
0	63903	7250
1	170	409

Model – Gradient Boosting Machine

- Random splitting
 - Logistic Regression:
 - Accuracy train: **93.9 %**

prediction	0	1
0	85771	5072
1	643	2265

- Accuracy test: **93.4%**

prediction	0	1
0	46142	2969
1	359	1012

- Hyperparameters:
 - 10-fold Cross-validation

- nr trees: 800
- depth= 2
- shrinkage =0.01

In a binary setting:

		Predicted	
		P	N
Condition	P	TP	FN
	N	FP	TN

- $\text{Recall} = \text{TP} / (\text{TP} + \text{FN}) = 46142 / (46142 + 359) = 46142 / 46501 = \mathbf{99.2 \%}$
- $\text{Precision} = \text{TP} / (\text{TP} + \text{FP}) = 46142 / (46142 + 2969) = 46142 / 49111 = \mathbf{93.9 \%}$
- $\text{True negative rate} = \text{TN} / (\text{TN} + \text{FP}) = 1012 / (1012 + 2969) = 1012 / 3981 = \mathbf{25.4 \%}$
- $\text{Balanced accuracy} = \text{Recall} + \text{TNR} / 2 = \mathbf{62 \%}$
- **Still better than random guess!**

Conclusions

- ▶ We had a quite high sparse data set
- ▶ Data reduction analysis:
 - ▶ PCA analysis
 - ▶ 270 numerical values allowed reduction to 20 variables that explain 80 % of the variance,
 - ▶ Average was input in the empty instances
 - ▶ Random Forests
 - ▶ Allowed discriminate the most relevant categorical variables for this classification problem
 - ▶ Variables with more than 53 factors were removed
 - ▶ Most of the remaining categorical variables were relevant except one
- ▶ Gradient Boosting Machine with the selected 38 variables:
 - ▶ Apparently not a time series problem ("reported chargeback on the card as fraud transaction ")
 - ▶ 93.4 % accuracy
 - ▶ **In future use the balanced accuracy for optimizing the hyperparameters**